

One-Hot Encoding and Trainable Word Embeddings

Overview

One-Hot Encoding -

- Creates a binary vector for each category where only one element is "hot" (1) and others are 0
- Dimensionality equals the number of unique categories
- No inherent relationship between categories
- Fixed representation that doesn't change during training

Trainable Word Embeddings -

- Learns dense, low-dimensional vector representations during training
- Dimensionality is a hyperparameter (typically much smaller than one-hot)
- Captures relationships between categories through learned similarity
- Representation improves as the model trains

Performance Comparison

Advantages of One-Hot Encoding -

Simplicity: Easy to implement and understand

No information loss: Preserves all categorical distinctions

Works well with small categorical spaces

Deterministic: Doesn't require learning

Advantages of Trainable Word Embeddings -

Dimensionality reduction: More efficient for large vocabularies

Learned semantics: Can capture meaningful relationships

Better generalization: Often improves model performance

Memory efficiency: Requires less space for large categorical variables

Observations from Text Generation Experiment

```
D:\Dulya>python text_generation.py
Vocabulary Size: 4752
Training One-Hot Encoding Model...
Epoch 1, Loss: 6.5587
Epoch 2, Loss: 5.9977
Epoch 3, Loss: 5.6373
Epoch 4, Loss: 5.2530
Epoch 5, Loss: 4.8426
Epoch 6, Loss: 4.4019
Epoch 7, Loss: 3.9439
Epoch 8, Loss: 3.4774
Epoch 9, Loss: 3.0109
Epoch 10, Loss: 2.5587

Training Trainable Embeddings Model...
Epoch 1, Loss: 6.3774
Epoch 2, Loss: 5.0130
Epoch 3, Loss: 5.1509
Epoch 4, Loss: 4.7061
Epoch 5, Loss: 4.2640
Epoch 6, Loss: 3.8257
Epoch 7, Loss: 3.3947
Epoch 8, Loss: 2.9854
Epoch 9, Loss: 2.6029
Epoch 10, Loss: 2.2502

One-Hot Encoding: Training Time: 510.68 seconds, Final Loss: 2.5587
Trainable Embeddings: Training Time: 112.52 seconds, Final Loss: 2.2502

Generated Text with One-Hot Encoding:
o my luve ' s like the earth , that may be in a man , and i tell it , i know in the

Generated Text with Trainable Embeddings:
o my luve ' s the sweep of easy wind and downy flake . the little plentiful manikins skipping around in collars and tail '

Quick Comparison:
- One-Hot: Simple but uses more memory, may not capture word meanings well.
- Embeddings: Faster, learns word relationships, likely generates better text.

D:\Dulya>
```

1. Training Performance -

Convergence Speed :

Trainable embeddings achieved lower loss values consistently across all epochs (6.27 → 2.25 vs 6.55 → 2.55)

Faster convergence suggests embeddings learn more efficient representations

Training Time :

Embeddings were **4.5x faster** (112s vs 510s)

Significant computational advantage despite same epoch count

Final Loss :

Embeddings achieved better final loss (2.25 vs 2.55)

12% relative improvement in optimization

2. Generated Text Quality -

One-Hot Output :

Shows basic syntactic correctness but limited semantic coherence

Repetitive structure ("i tell it, i know in the")

Lacks contextual flow between phrases

Embeddings Output :

Displays richer poetic imagery ("sweep of easy wind and downy flake")

Better lexical diversity ("plentiful manikins skipping around")

More creative word combinations suggesting learned semantic relationships

3. Technical Tradeoffs -

Memory Efficiency :

One-hot encoding with 4,752 vocabulary size creates 4,752-dim sparse vectors

Embeddings likely used \ll 100 dimensions (typical range 50-300)

Representation Power :

One-hot treats all words as equally distant

Embeddings naturally cluster semantically similar words

Key Findings

- For text generation, embeddings clearly outperform one-hot encoding in :
 - Training efficiency (time and convergence)
 - Output quality (coherence and creativity)
 - Memory usage (despite larger raw vocabulary)
- Surprising Result :

The performance gap is larger than expected - embeddings are both faster AND better, contrary to the common assumption that better quality requires more computation

Recommendations

- Always prefer embeddings for NLP tasks with vocabulary sizes >1000
- The observed 4.5x speedup justifies using embeddings even for prototyping
- For production systems, the memory savings become increasingly valuable at scale

Limitations

- Without perplexity/metrics, qualitative assessment is subjective
- Hardware differences could affect timing results
- Small sample size of generated text (single example each)